

CSE 489/589

Programming Assignment 3

Outline

Part 1 – Assignment Overview

- Problem statement
- Programming environment
- Topology
- Differences from Text-Book ...

Part 2 – Socket Programming

- UDP socket overview
- select(), again!
- Packet Formats

Part 1:

Assignment Overview

Problem statement

- Implement a simplified version of the Distance Vector Protocol
- Basic idea:
 - When x receives new DV estimate from neighbor, it updates its own DV using Bellman-Ford equation:

$$D_x(y) \leftarrow \min_v \{c(x, v) + D_v(y)\} \quad \text{for each node } y \in N$$

Distance vector algorithm

Bellman-Ford equation (dynamic programming)

let

$d_x(y) := \text{cost of least-cost path from } x \text{ to } y$

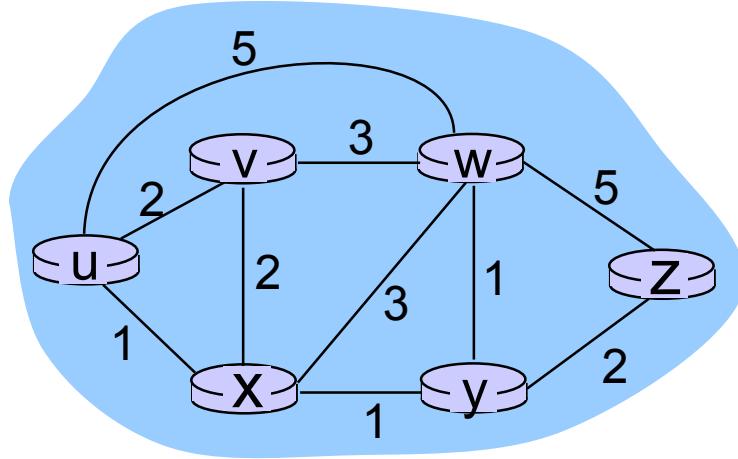
then

$$d_x(y) = \min_v \{ c(x, v) + d_v(y) \}$$

| | |
| | |
 v cost from neighbor v to destination y
cost to neighbor v

 \min taken over all neighbors v of x

Bellman-Ford example



clearly, $d_v(z) = 5$, $d_x(z) = 3$, $d_w(z) = 3$

B-F equation says:

$$\begin{aligned}d_u(z) &= \min \{ c(u,v) + d_v(z), \\&\quad c(u,x) + d_x(z), \\&\quad c(u,w) + d_w(z) \} \\&= \min \{ 2 + 5, \\&\quad 1 + 3, \\&\quad 5 + 3 \} = 4\end{aligned}$$

node achieving minimum is next
hop in shortest path, used in forwarding table

Distance vector algorithm

- $D_x(y)$ = estimate of least cost from x to y
 - x maintains distance vector $\mathbf{D}_x = [D_x(y): y \in N]$
- node x :
 - knows cost to each neighbor v : $c(x,v)$
 - maintains its neighbors' distance vectors. For each neighbor v , x maintains
 $\mathbf{D}_v = [D_v(y): y \in N]$

Distance vector algorithm

key idea:

- from time-to-time, each node sends its own distance vector estimate to neighbors
- when x receives new DV estimate from neighbor, it updates its own DV using B-F equation:

$$D_x(y) \leftarrow \min_v \{c(x,v) + D_v(y)\} \text{ for each node } y \in N$$

- ❖ under minor, natural conditions, the estimate $D_x(y)$ converge to the actual least cost $d_x(y)$

Distance vector algorithm

iterative, asynchronous:

each local iteration
caused by:

- local link cost change
- DV update message from neighbor

distributed:

- each node notifies neighbors *only* when its DV changes
 - neighbors then notify their neighbors if necessary

each node:

wait for (change in local link cost or msg from neighbor)

recompute estimates

if DV to any dest has changed, *notify* neighbors

$$D_x(y) = \min\{c(x,y) + D_y(y), c(x,z) + D_z(y)\}$$

$$= \min\{2+0, 7+1\} = 2$$

$$D_x(z) = \min\{c(x,y) + D_y(z), c(x,z) + D_z(z)\}$$

$$= \min\{2+1, 7+0\} = 3$$

**node x
table**

	cost to		
	x	y	z
x	0	2	7
y	∞	∞	∞
z	∞	∞	∞

**node y
table**

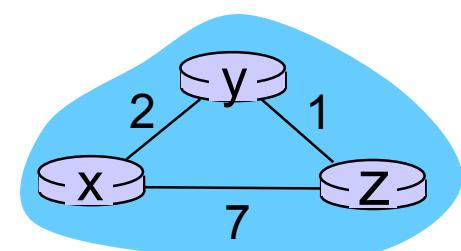
	cost to		
	x	y	z
x	∞	∞	∞
y	2	0	1
z	∞	∞	∞

**node z
table**

	cost to		
	x	y	z
x	∞	∞	∞
y	∞	∞	∞
z	7	1	0

	cost to		
	x	y	z
x	0	2	3
y	2	0	1
z	7	1	0

time



$$D_x(y) = \min\{c(x,y) + D_y(y), c(x,z) + D_z(y)\}$$

$$= \min\{2+0, 7+1\} = 2$$

$$D_x(z) = \min\{c(x,y) + D_y(z), c(x,z) + D_z(z)\}$$

$$= \min\{2+1, 7+0\} = 3$$

**node x
table**

	cost to		
	x	y	z
from	x	0 2 7	
y	∞ ∞ ∞		
z	∞ ∞ ∞		

**node y
table**

	cost to		
	x	y	z
from	x	∞ ∞ ∞	
y	2 0 1		
z	∞ ∞ ∞		

**node z
table**

	cost to		
	x	y	z
from	x	∞ ∞ ∞	
y	∞ ∞ ∞		
z	7 1 0		

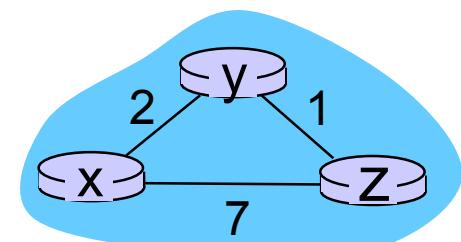
	cost to		
	x	y	z
from	x	0 2 3	
y	2 0 1		
z	7 1 0		

	cost to		
	x	y	z
from	x	0 2 3	
y	2 0 1		
z	3 1 0		

	cost to		
	x	y	z
from	x	0 2 3	
y	2 0 1		
z	3 1 0		

	cost to		
	x	y	z
from	x	0 2 3	
y	2 0 1		
z	3 1 0		

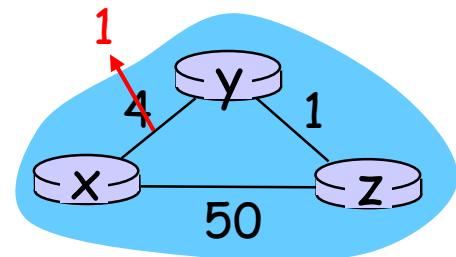
time



Distance vector: link cost changes

link cost changes:

- ❖ node detects local link cost change
- ❖ updates routing info, recalculates distance vector
- ❖ if DV changes, notify neighbors



**“good
news
travels
fast”**

t_0 : y detects link-cost change, updates its DV, informs its neighbors.

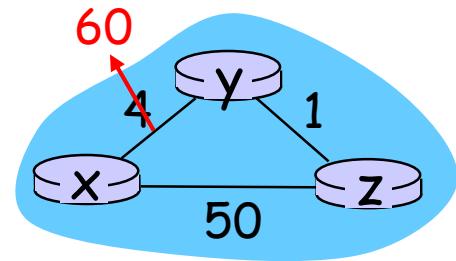
t_1 : z receives update from y, updates its table, computes new least cost to x , sends its neighbors its DV.

t_2 : y receives z' s update, updates its distance table. y' s least costs do *not* change, so y does *not* send a message to z.

Distance vector: link cost changes

link cost changes:

- ❖ node detects local link cost change
- ❖ *bad news travels slow* - “count to infinity” problem!
- ❖ 44 iterations before algorithm stabilizes: see text



poisoned reverse:

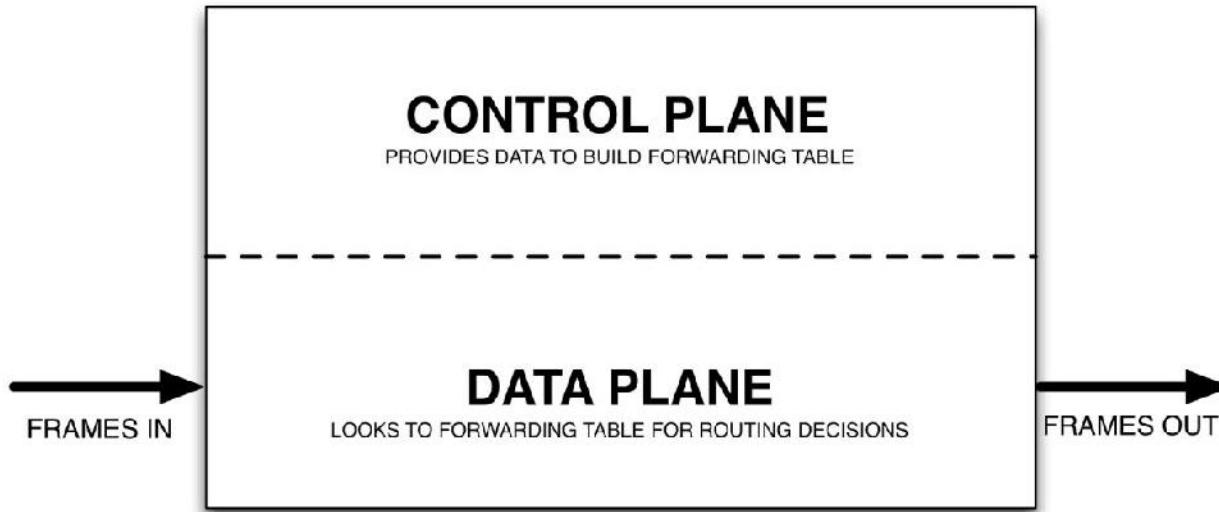
- ❖ If Z routes through Y to get to X :
 - Z tells Y its (Z's) distance to X is infinite (so Y won't route to X via Z)
- ❖ will this completely solve count to infinity problem?

Programming environment

- Use 5 CSE student servers as *routers*
 - stones
 - euston
 - highgate
 - embankment
 - underground

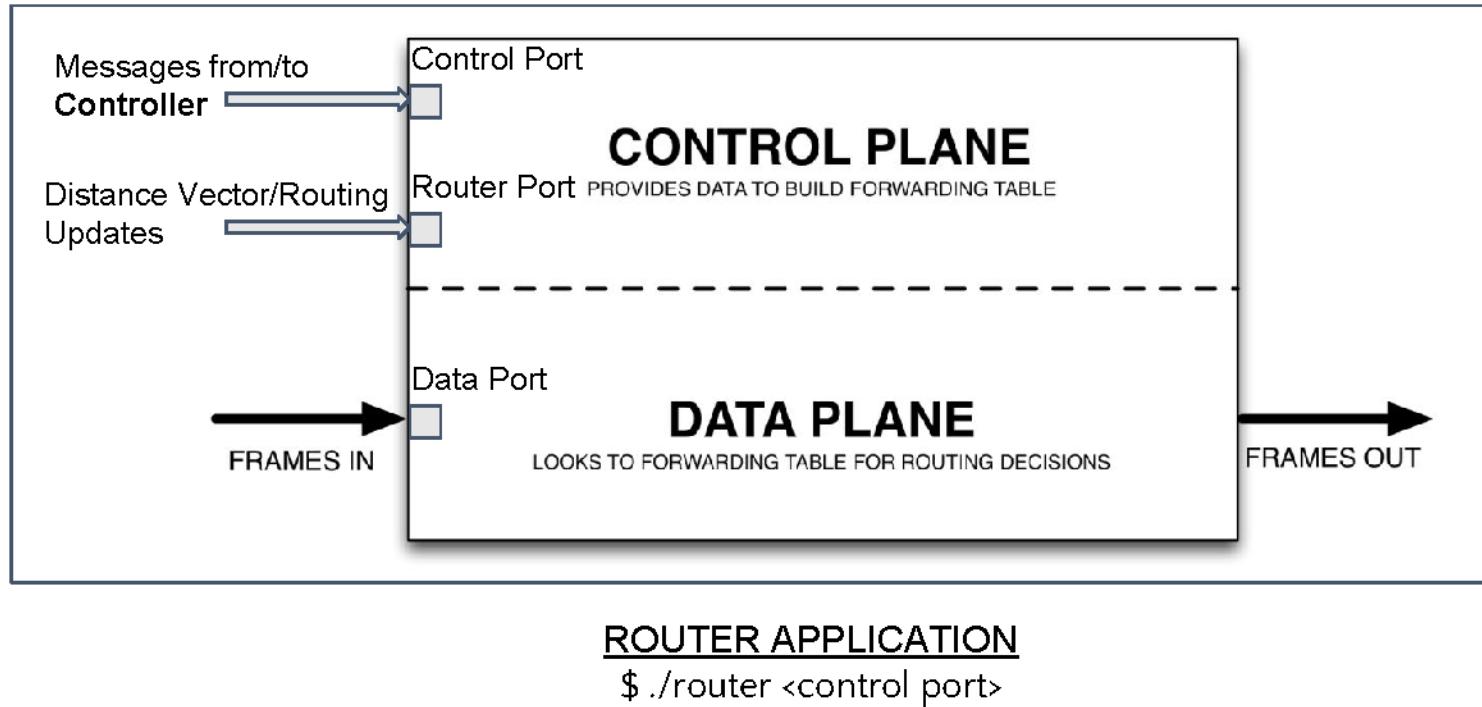
Routers

- Backbone of the internet



Routers: PA3

- Backbone of the internet



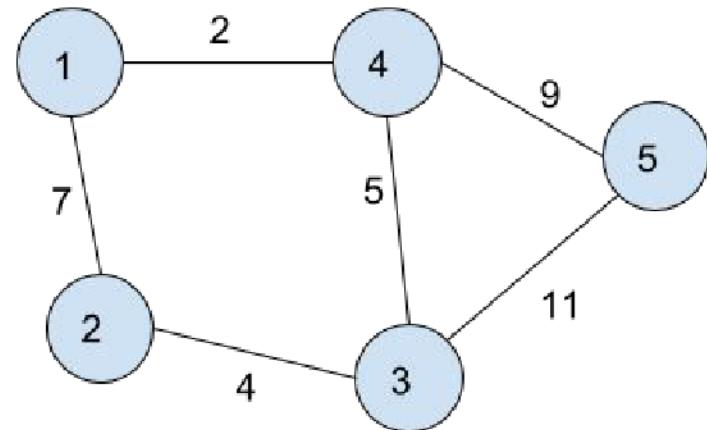
Controller

- A separate application
 - Generates *control* messages for routers
 - Expects response within a specified time limit
- Included in the template. No need to implement
- Router application needs to act and generate response packets to control messages

Topology

- Established by controller using INIT messages
- Reads a *topology file*
- Topology file
 - Snapshot of the network
 - Number of routers in the network and ID
 - IP address and all three port numbers
 - Links and their cost

Topology File: Example



Line #	Line entry
1	5
2	1128.205.36.8 4091 3452 2344
3	2128.205.35.24 4094 4562 2345
4	3128.205.36.24 4096 8356 1635
5	4128.205.36.4 7091 4573 1678
6	5128.205.36.25 7864 3456 1946
7	127
8	459
9	142
10	345
11	324
12	3511

<ID> <IP address> <control port> <router port> <data port>

<router ID 1> <router ID 2> <cost>

Differences from Text-Book ...

- Routing Updates: Periodic exchange Only
 - Not on DV changes
- To solve **count-to-infinity?**
 - No!

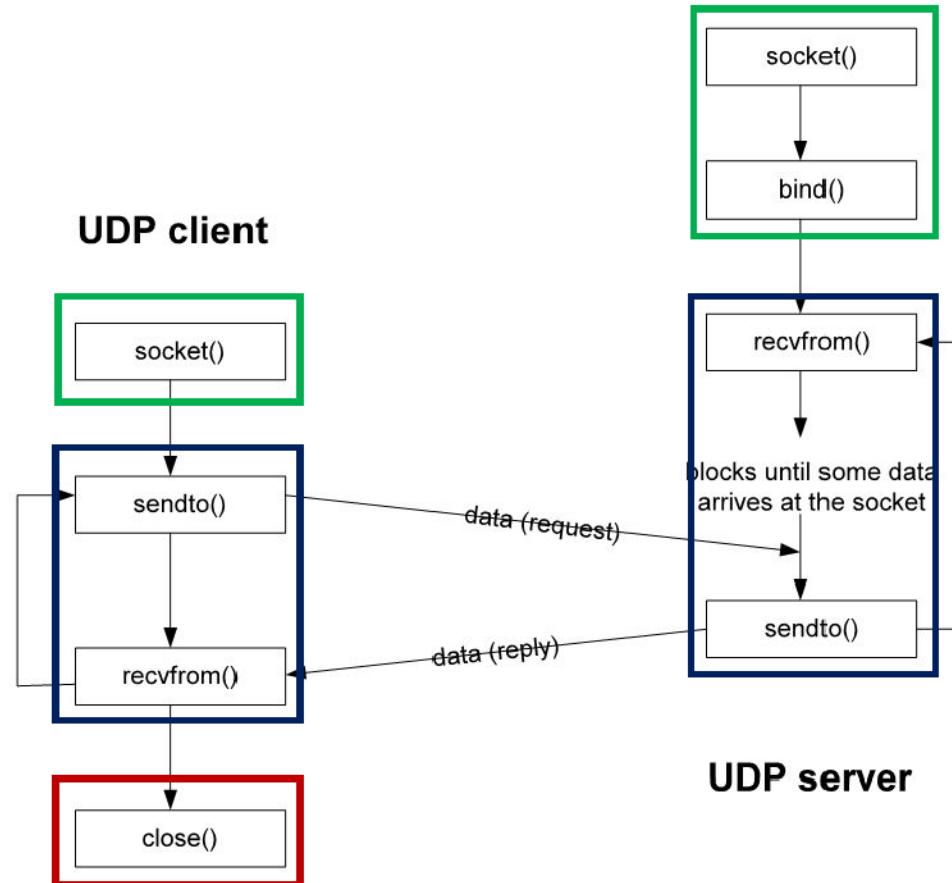
Part 2:

Socket Programming

UDP socket overview

- Client Init.(**socket**)
- Server Init. (**socket, bind**)
 - No connection establishment
- Data transfer
 - Client **sendto** - Server **recvfrom**
 - Server **sendto** - Client **recvfrom**

UDP socket overview



UDP server socket overview

```
server_socket = socket(AF_INET, SOCK_DGRAM, 0);
if(server_socket < 0)
    return err_msg_ERR("Cannot create socket");
```

```
bzero(&server_addr, sizeof(server_addr));

server_addr.sin_family = AF_INET;
server_addr.sin_addr.s_addr = htonl(INADDR_ANY);
server_addr.sin_port = htons(port);

if(bind(server_socket, (struct sockaddr *)&server_addr, sizeof(server_addr)) < 0 )
    return err_msg_ERR("Bind failed");
```

PA3: Multiplexing

- Control Port
 - **TCP**
 - Controller connections/messages
- Router Port
 - **UDP**
 - Distance vector updates
- Data Port
 - **TCP**
 - Actual data packets

select() again!

- I/O multiplexing with select
 - Control, router and data ports
 - Data on any connections established on these ports
- Timeout
 - To implement periodic distance vector broadcast

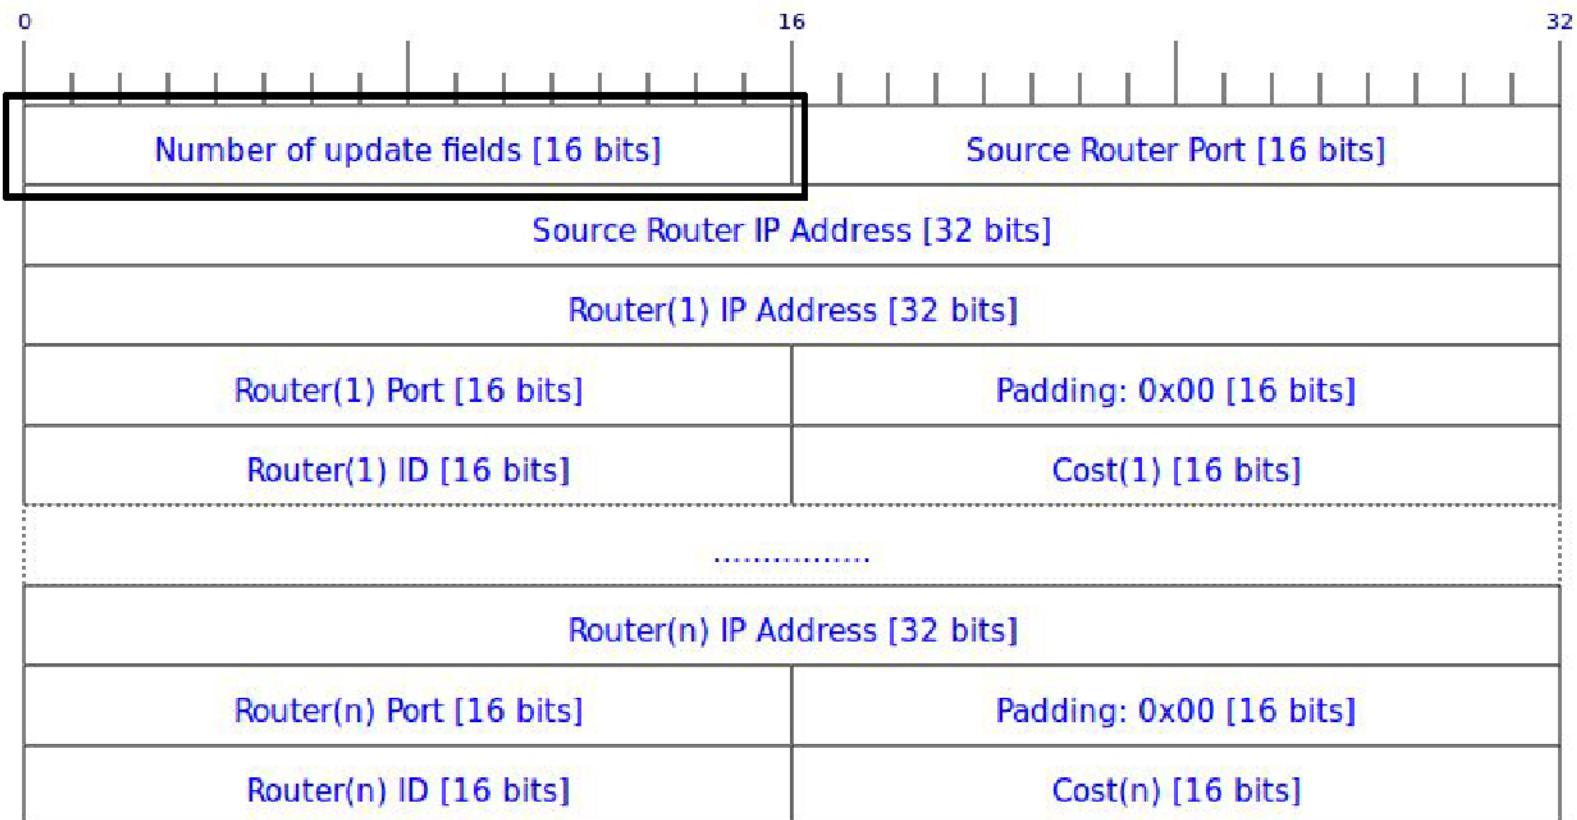
```
int select(int nfds, fd_set *readfds, fd_set *writefds, fd_set *exceptfds, struct timeval *timeout);
```

- Multiple timers needed!
 - Each router may be on a different schedule

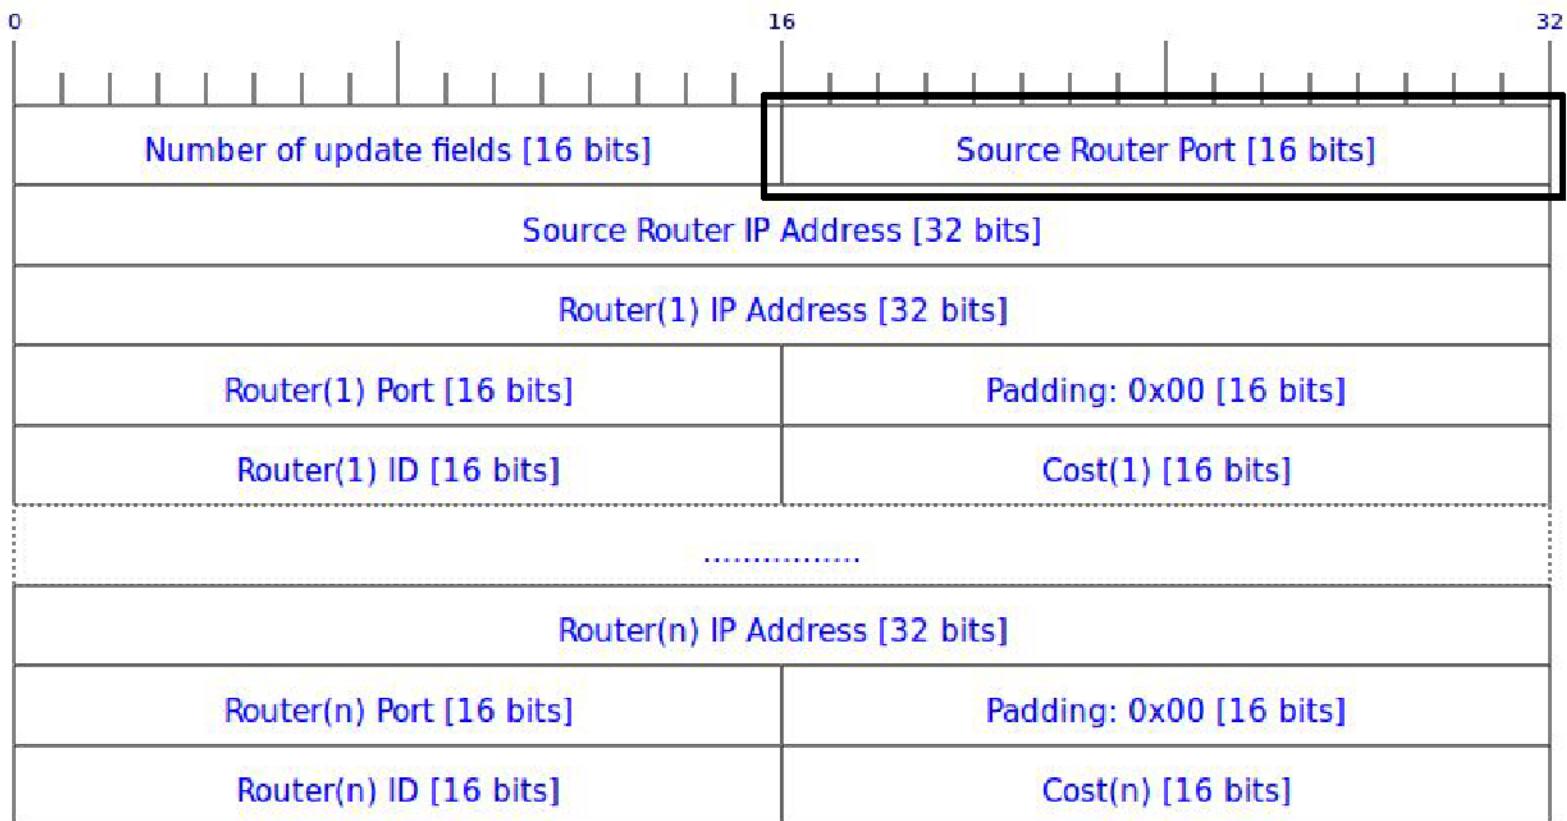
Packet Formats

- Multiple packet formats
 - Routing updates
 - Control
 - Control-Response
 - Data
- Need to be followed exactly!
- Example: Routing Update Packet

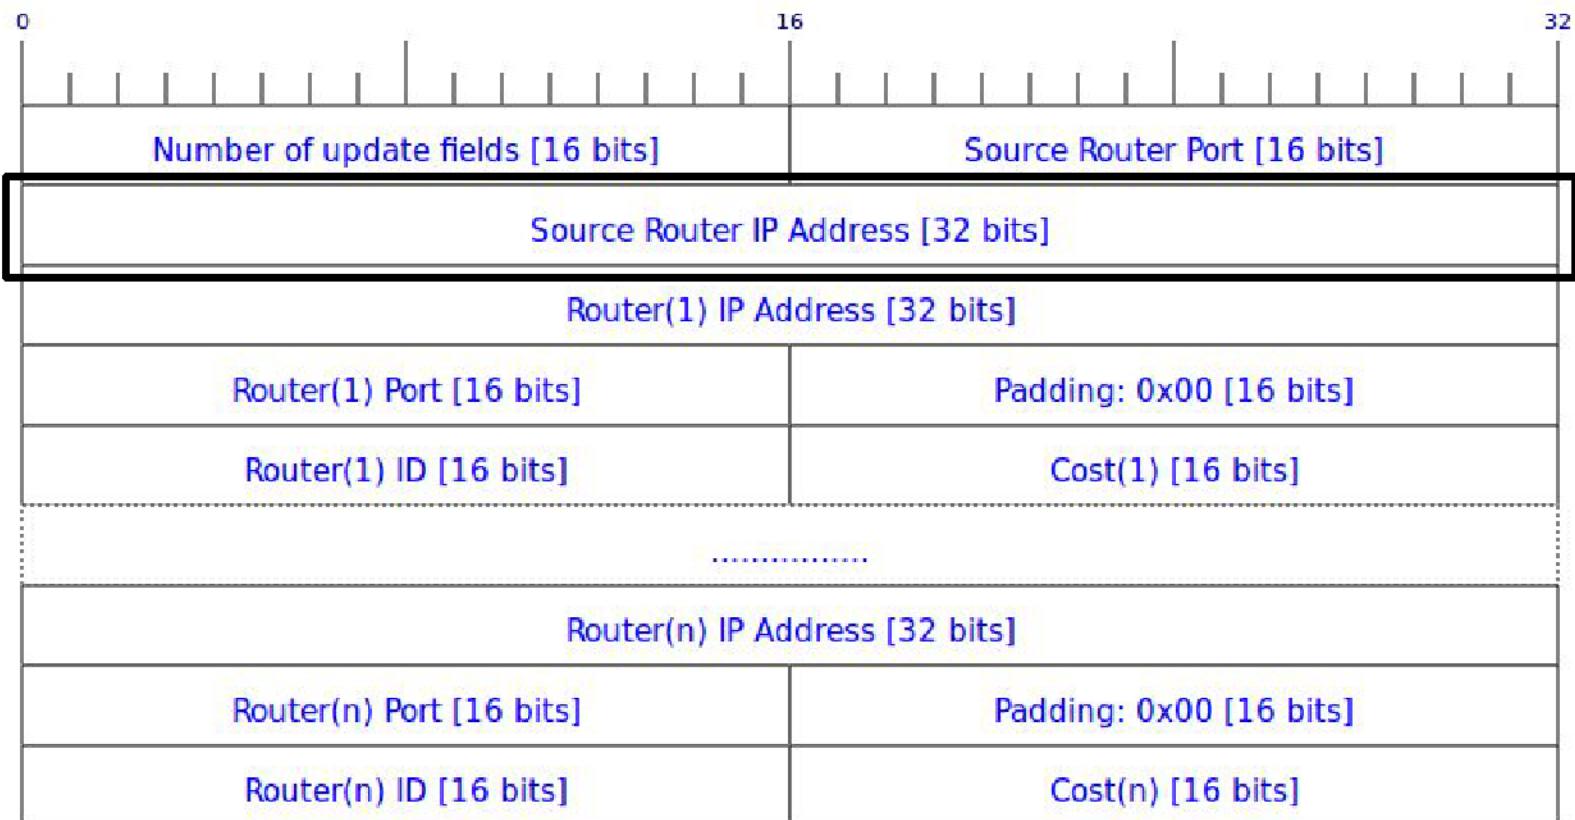
Packet Formats: Example



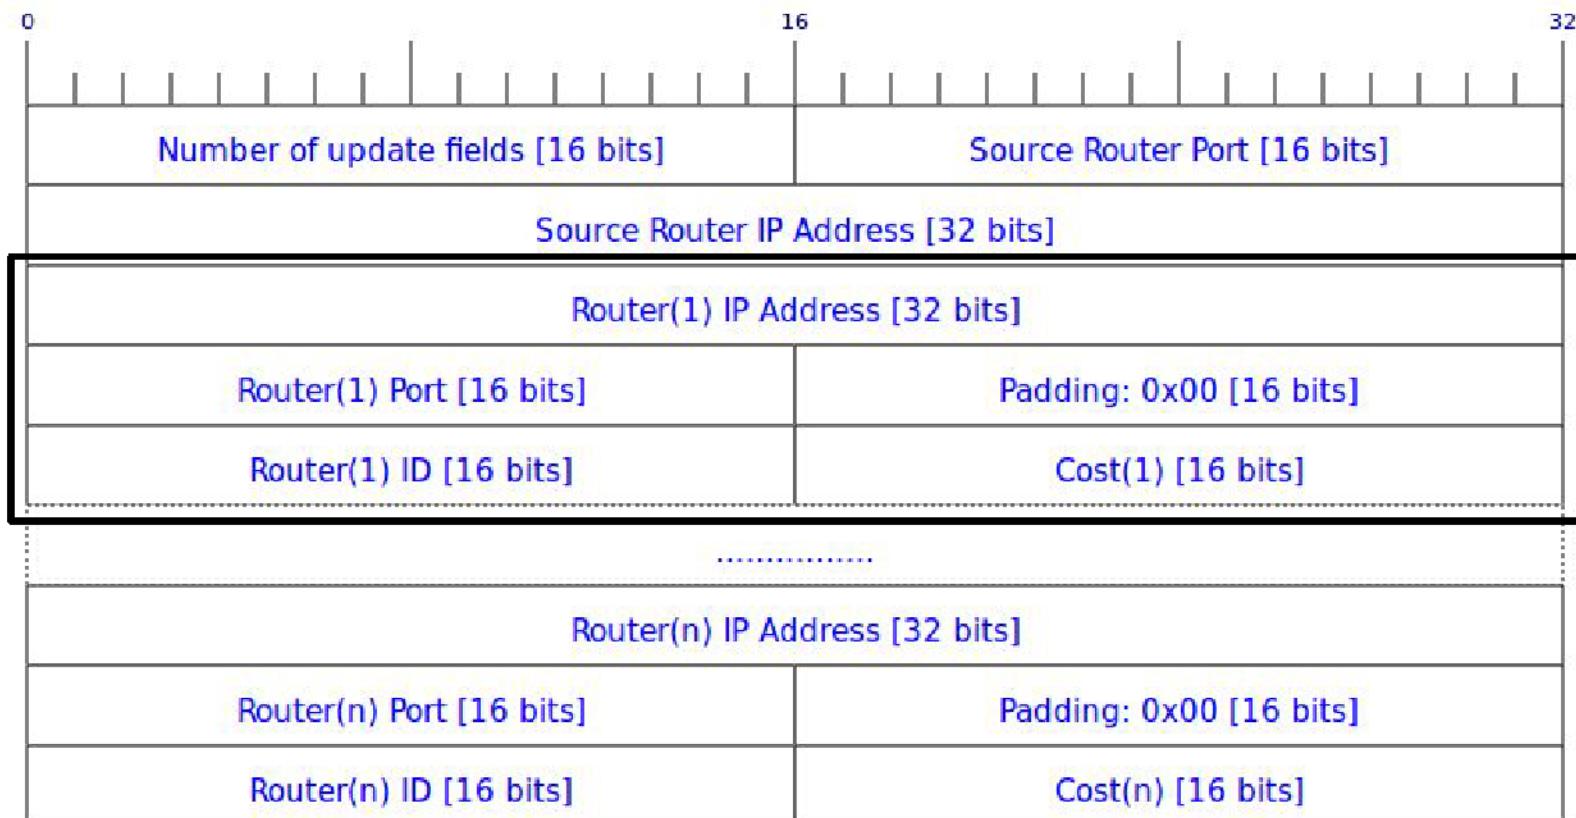
Packet Formats: Example



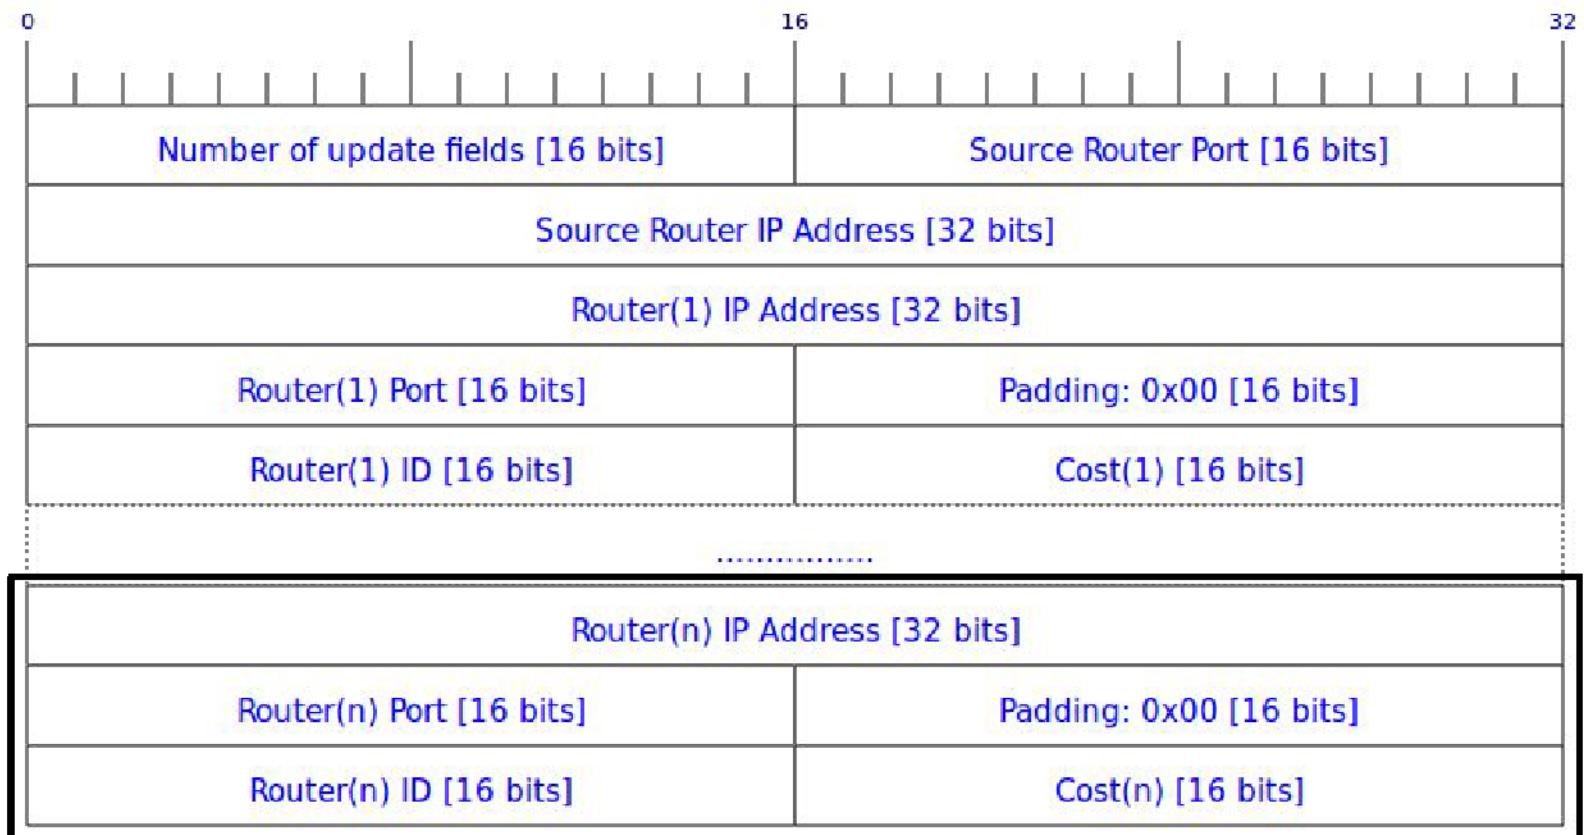
Packet Formats: Example



Packet Formats: Example



Packet Formats: Example



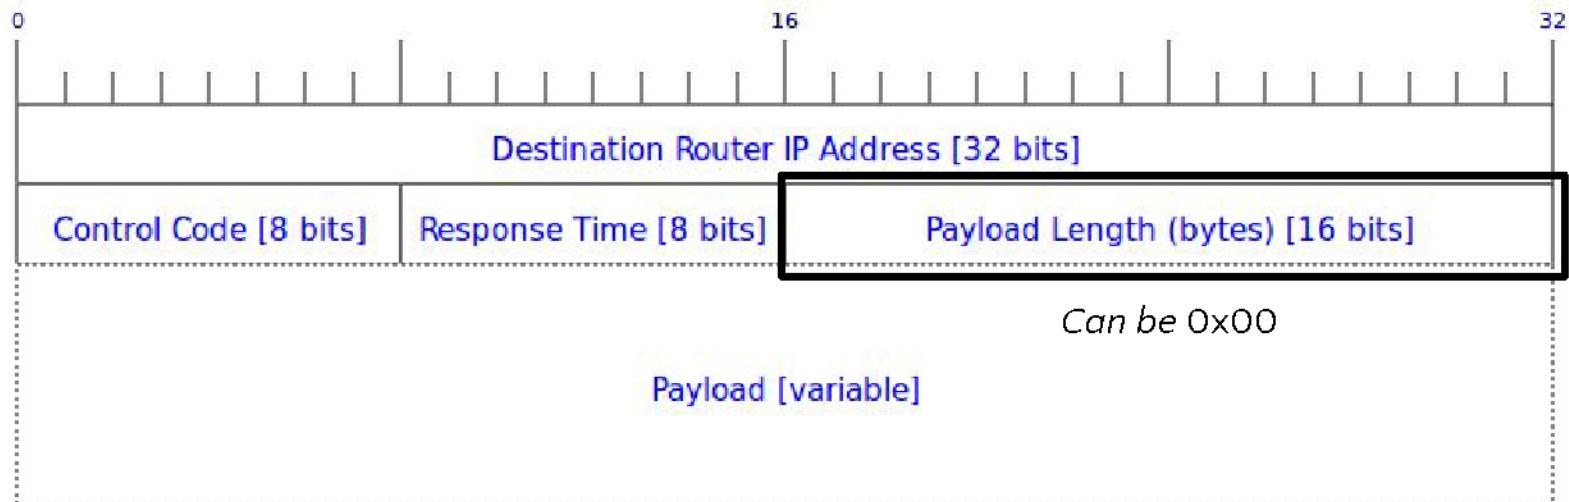
Numeric Types

- All numeric fields, unless mentioned otherwise, are represented/encoded as their **unsigned** versions
- How to represent infinity (INF)?
 - Largest 2-byte unsigned integer (65535)

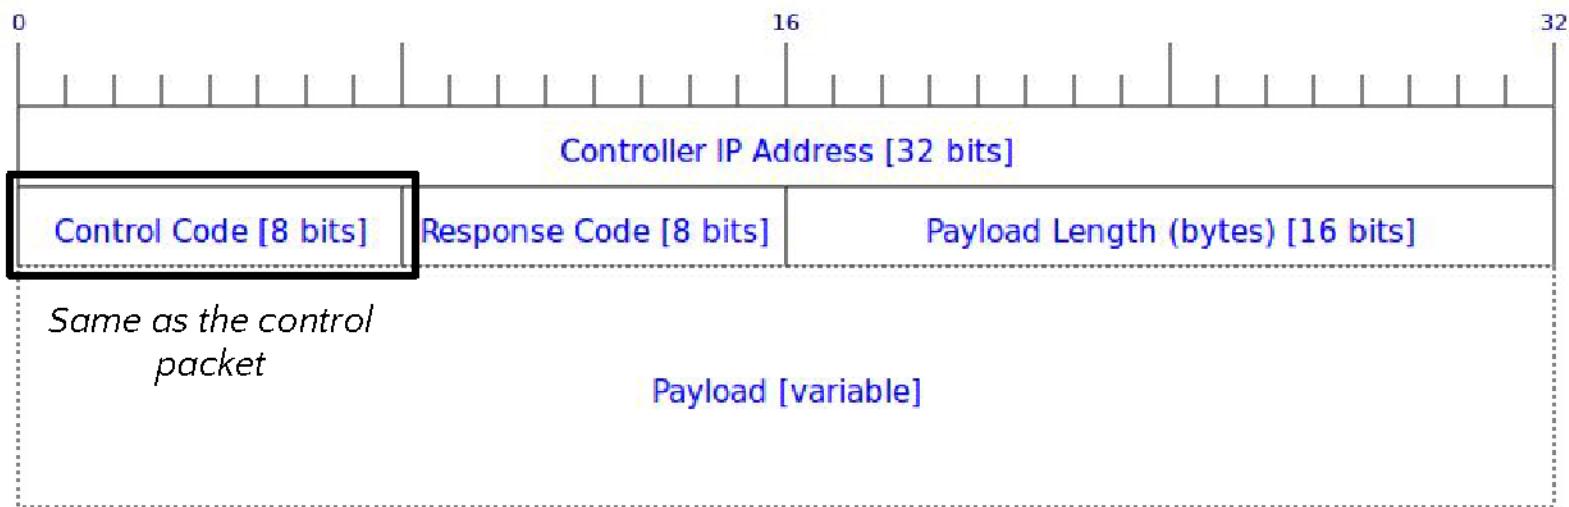
How do I verify my packets?

- Look at each bit manually!
- Can use tools like: hexdump
 - Byte-wise (hexadecimal)
\$ hexdump -C packet.file
 - Byte-wise (decimal)
 - Use the -e option with %u

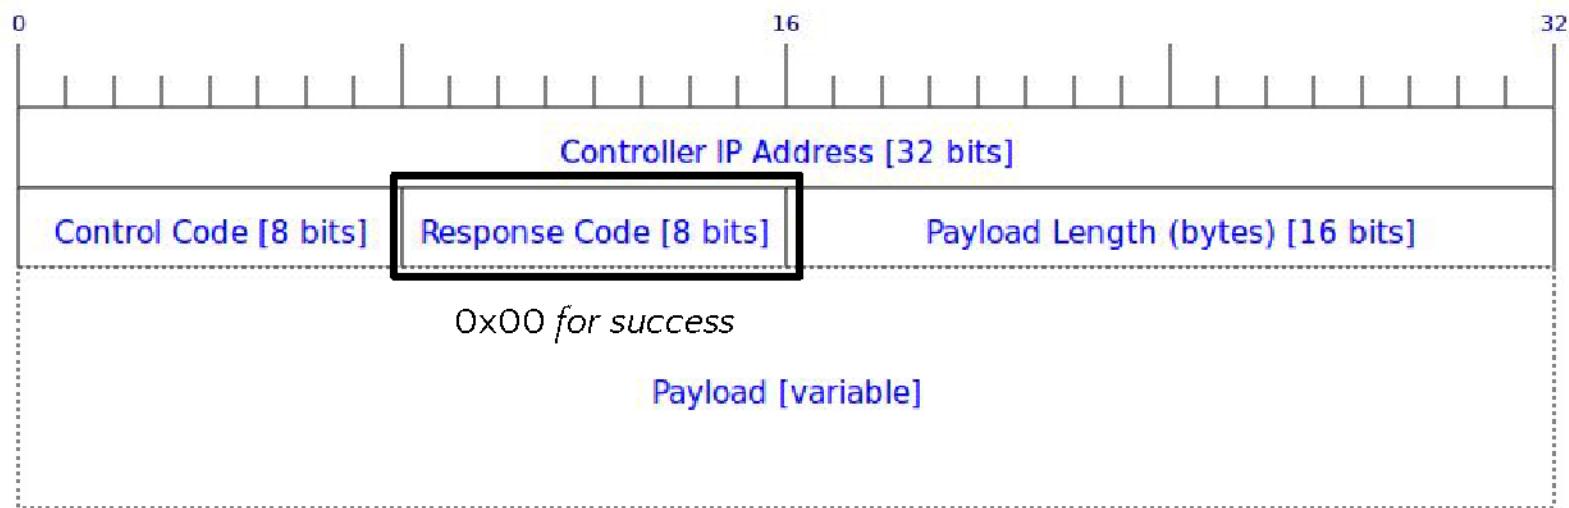
Control Header



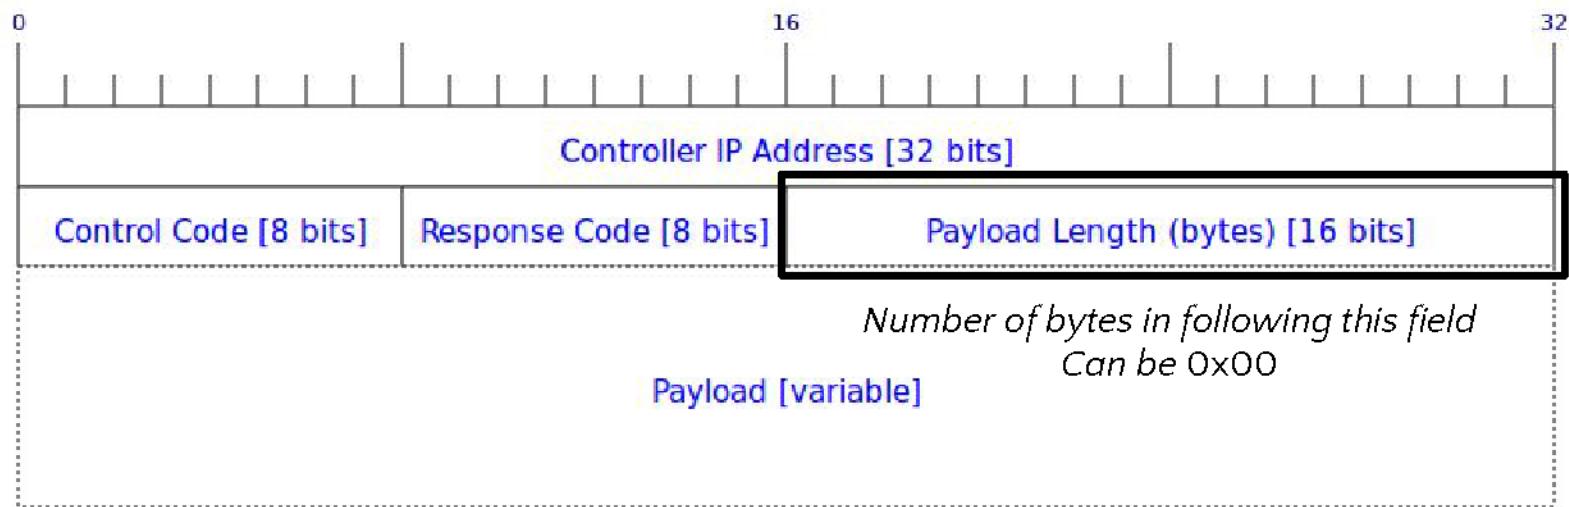
Control-Response Header



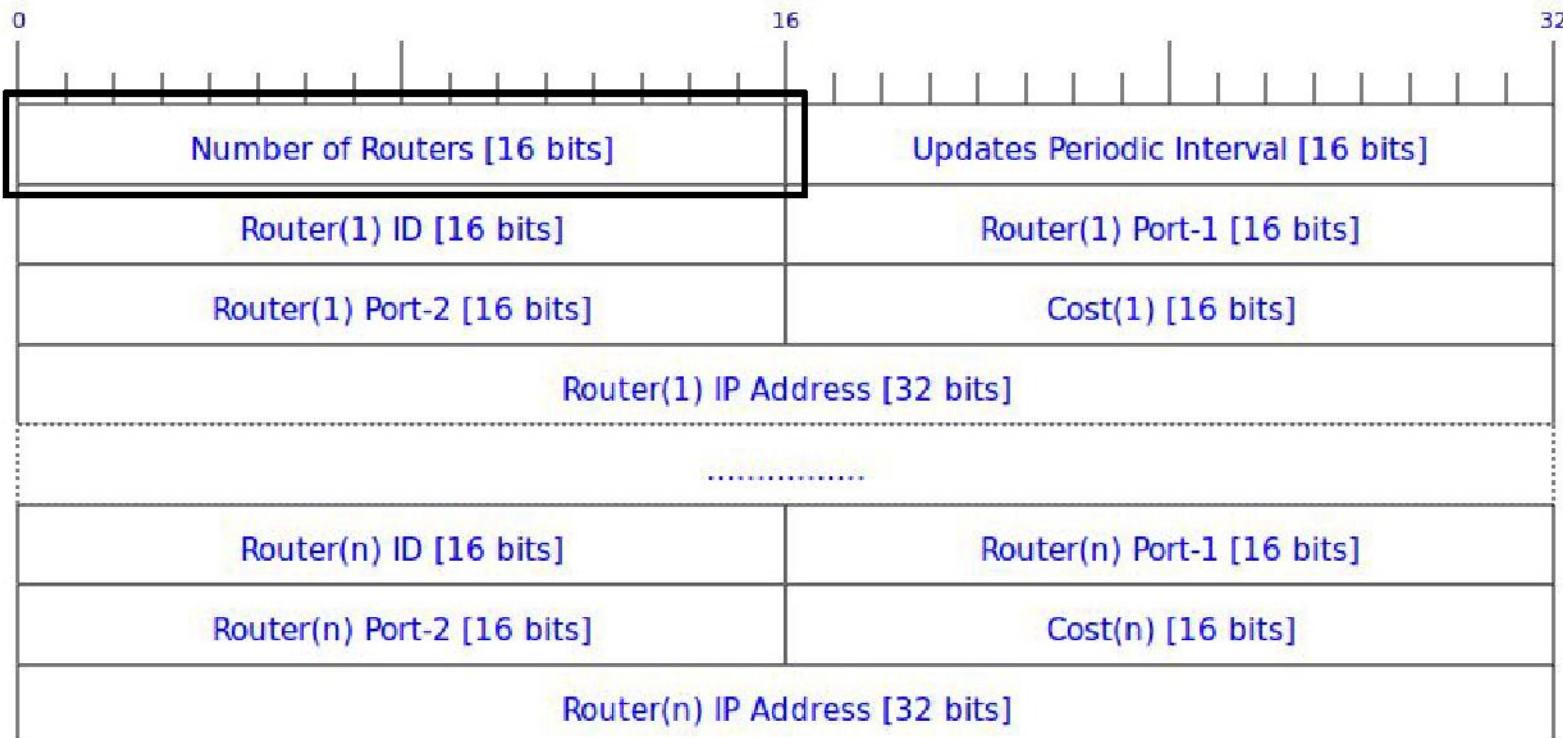
Control-Response Header



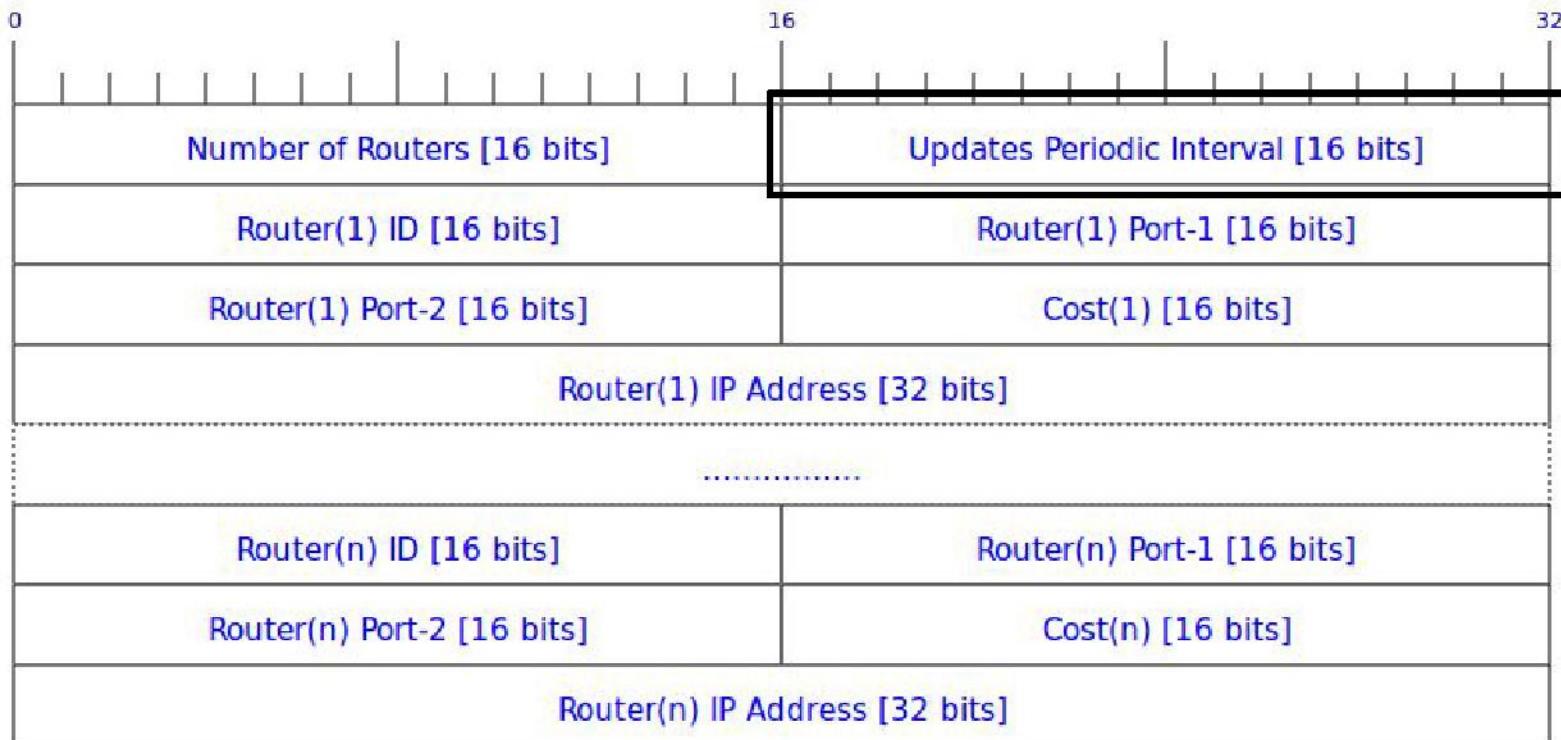
Control-Response Header



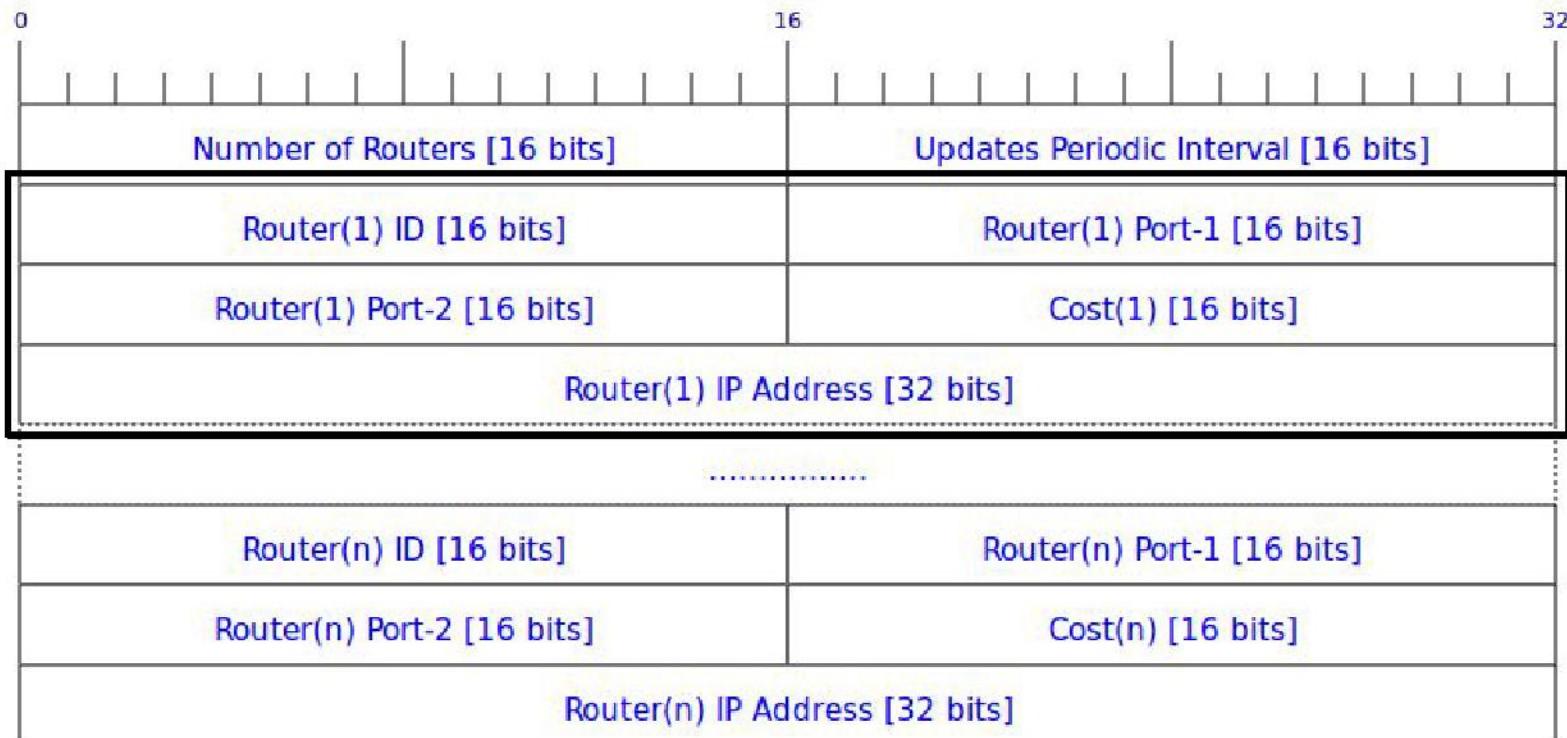
INIT (Control)



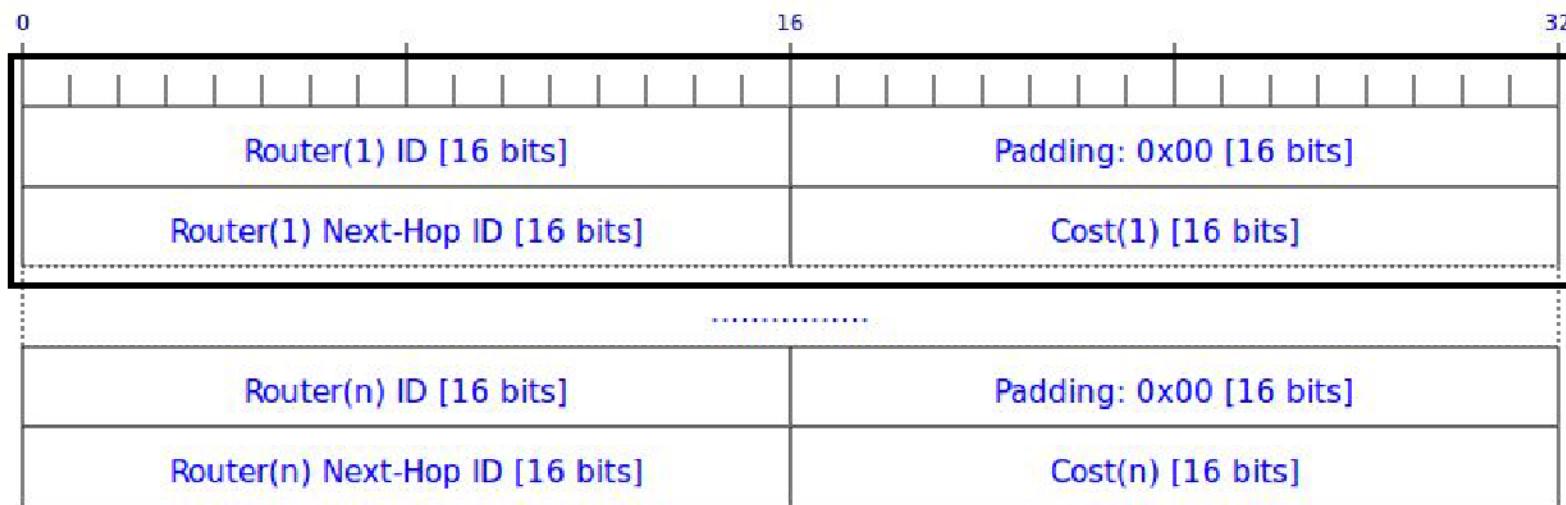
INIT (Control)



INIT (Control)

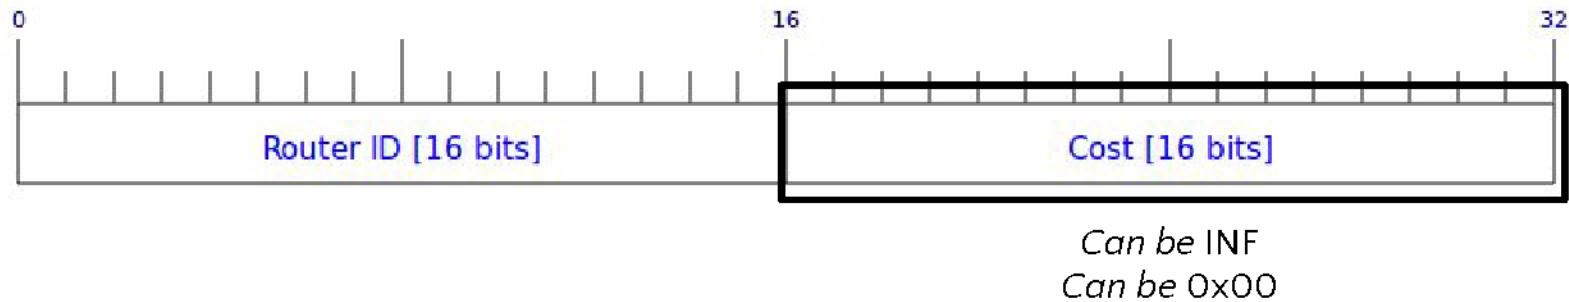


ROUTING-TABLE (Control-response)

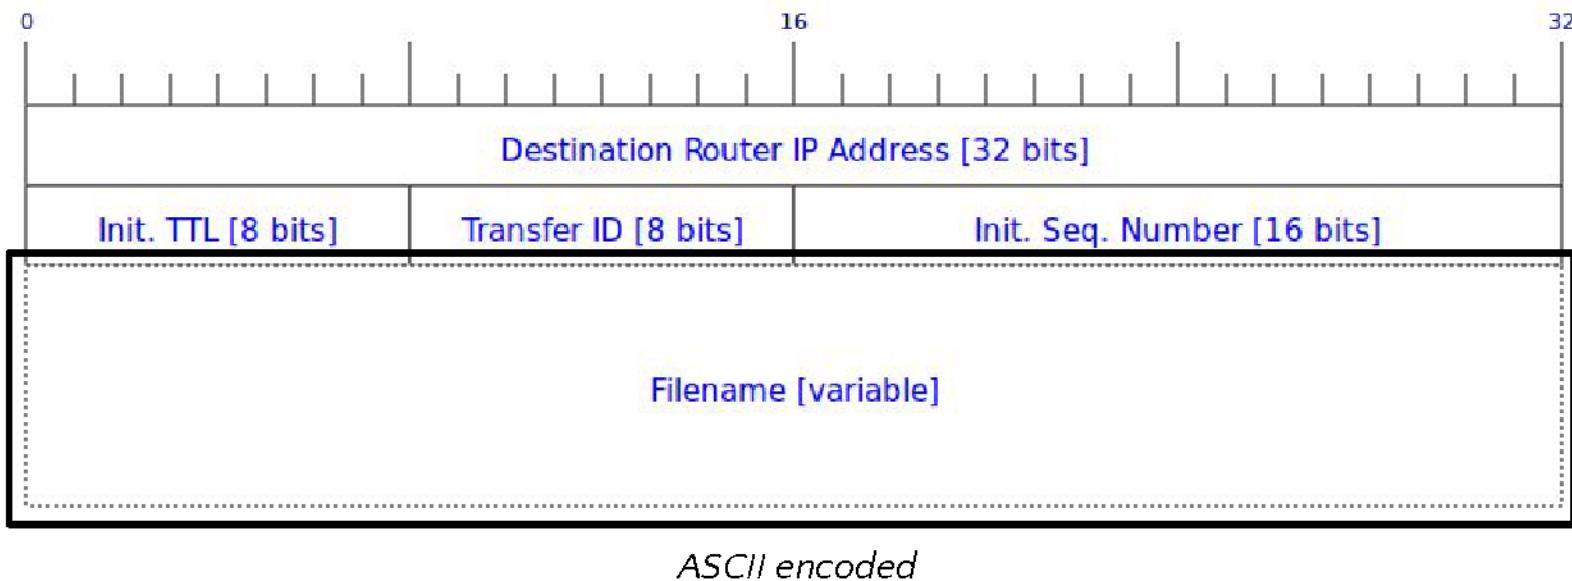


If you do not implement the ROUTING-TABLE response correctly, most automated tests will fail.

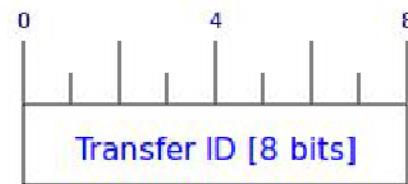
UPDATE (Control)



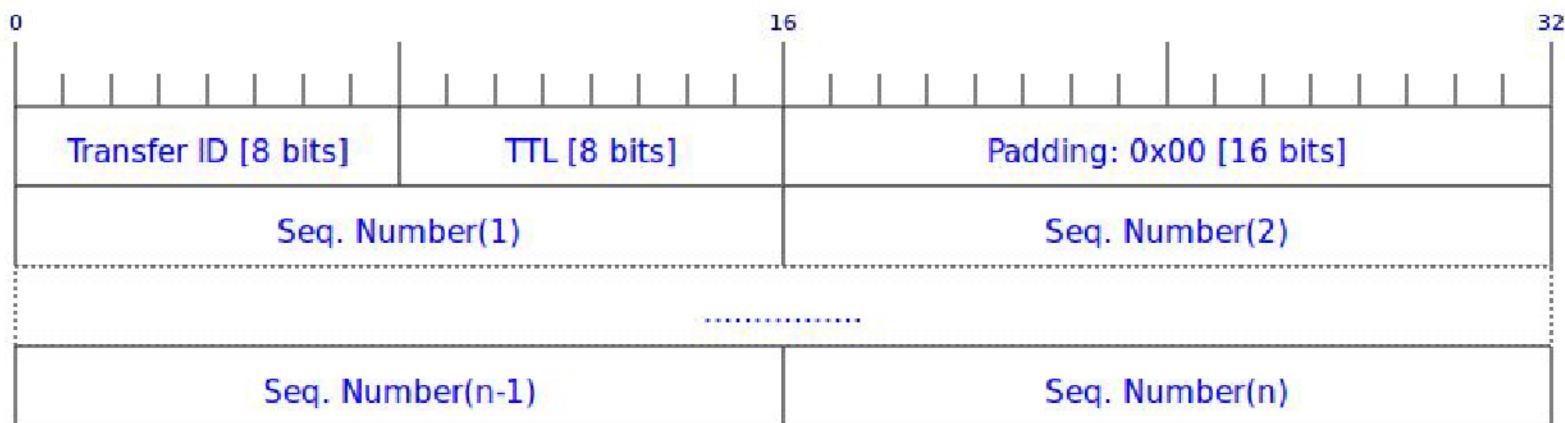
SENDFILE (Control)



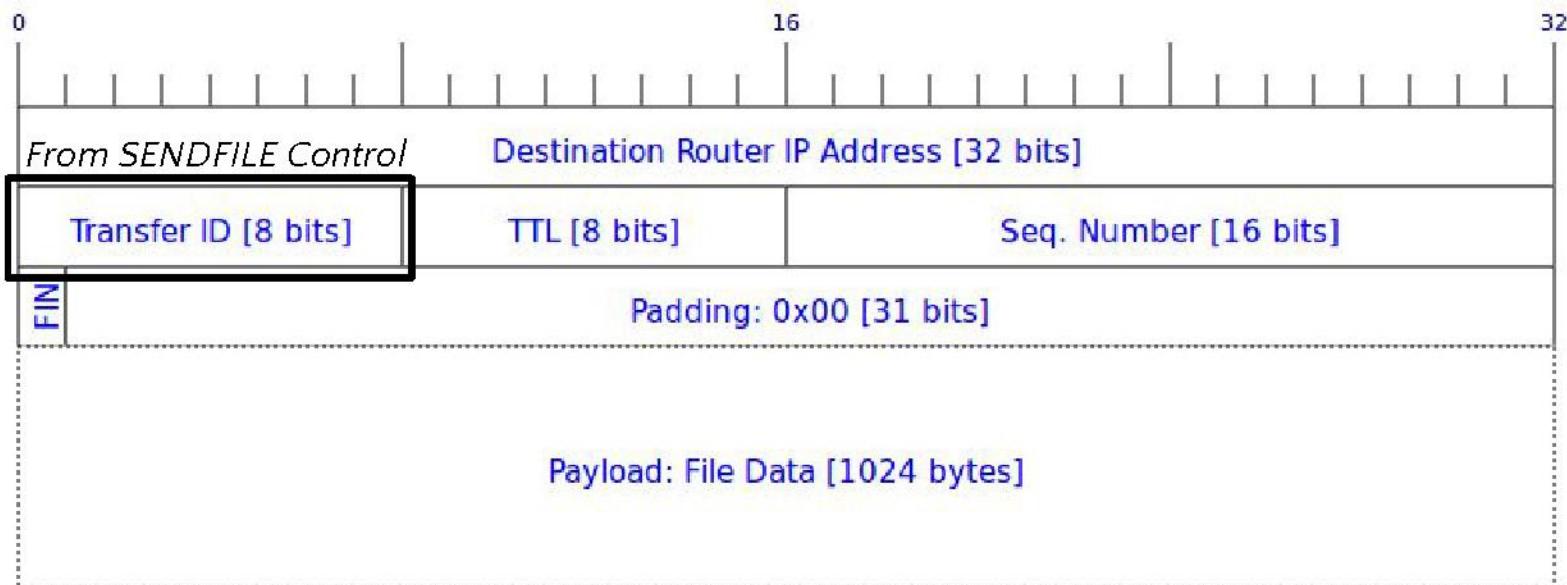
SENDFILE-STATS (Control)



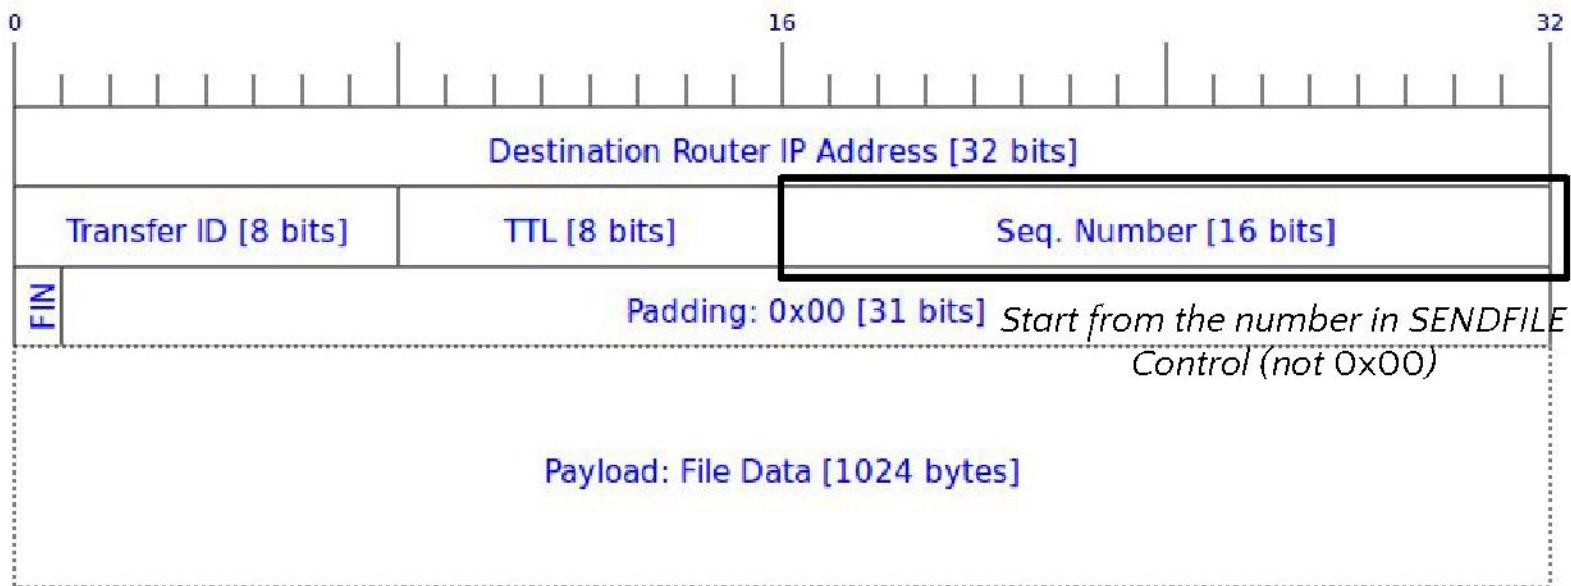
SENDFILE-STATS (Control-response)



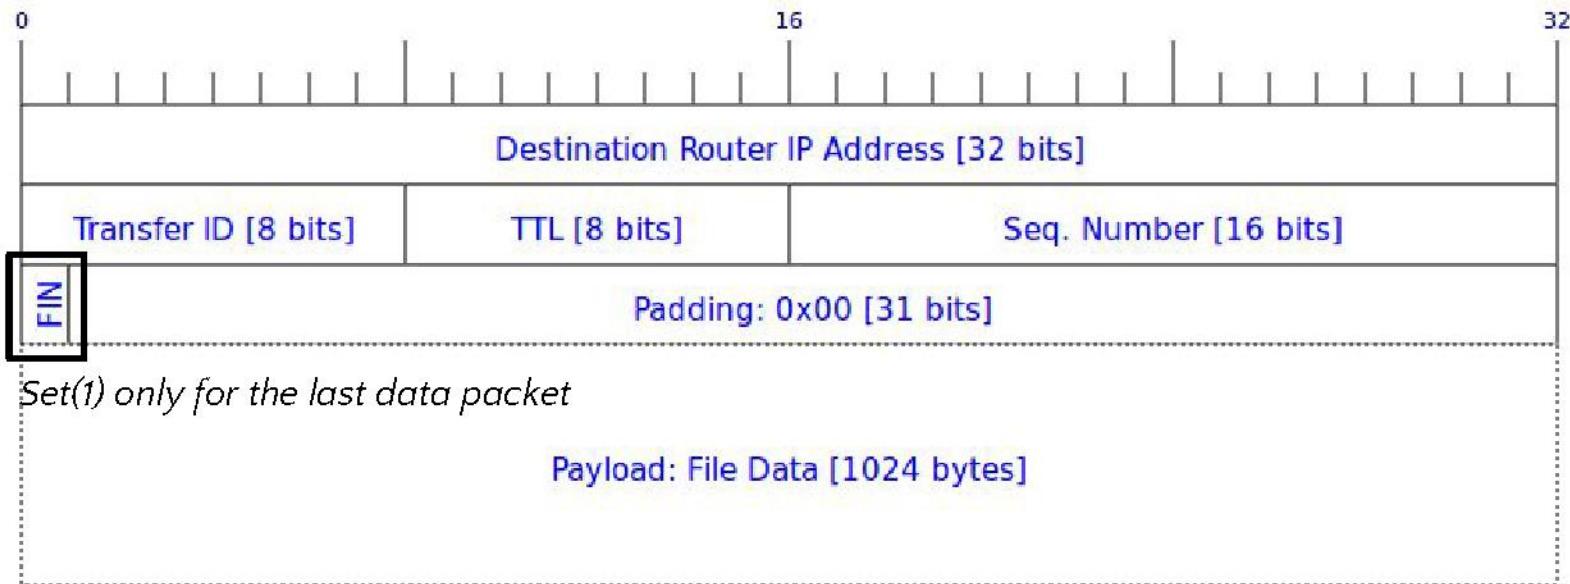
Data Packet



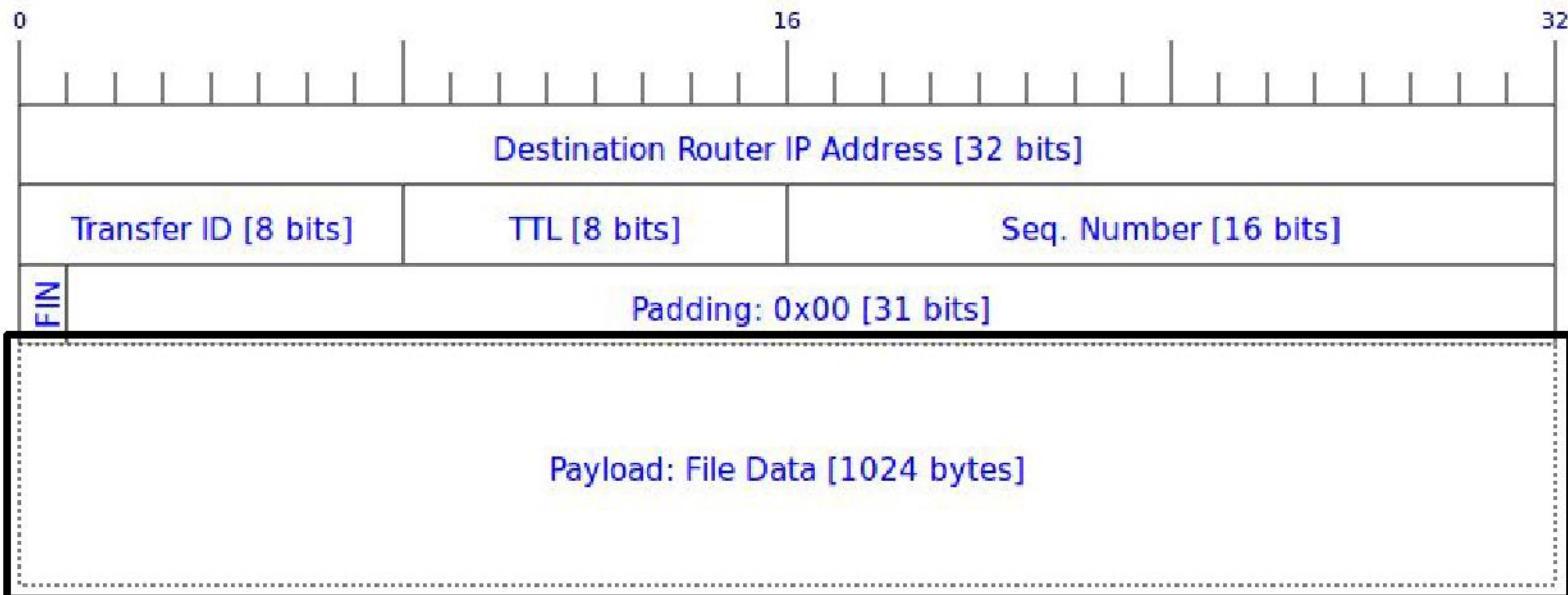
Data Packet



Data Packet



Data Packet



Should always be 1024 bytes. Assume filesize to be a multiple of 1024 bytes

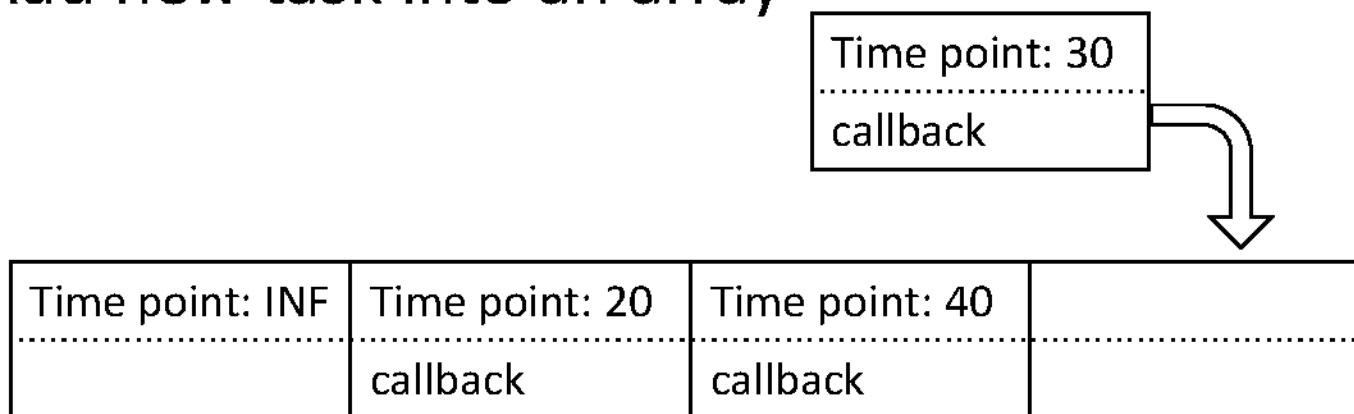
Timer: Basic idea

- Timeout
 - To implement periodic distance vector broadcast
 - To detect the absence of three consecutive routing updates
- Multiple timing tasks
 - Each router may be on a different schedule
- Single timer

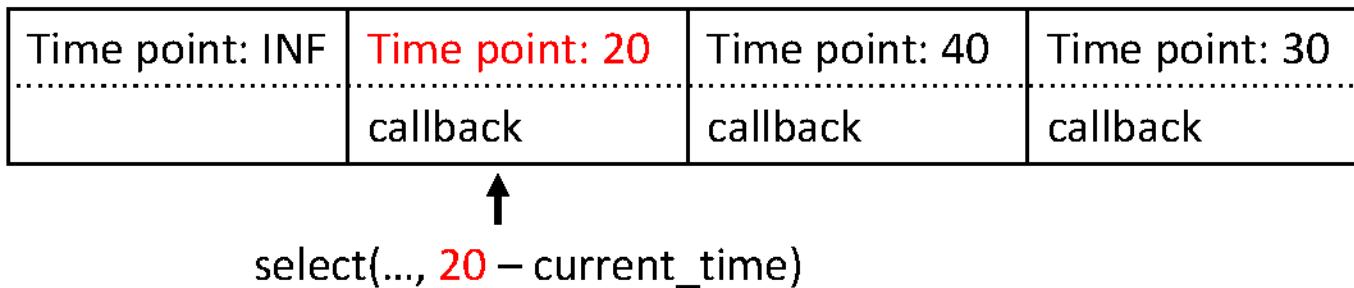
```
int select(int nfds, fd_set *readfds, fd_set *writefds, fd_set *exceptfds, struct timeval *timeout);
```

Timer: Start a timer

- Add new task into an array



- Time for the task with earliest starting time



Timer: Restart/Stop a timer

- Interruption & restart timer
 - Notice: select() can be interrupted by incoming messages as well as timeout!
 - You have to restart timer after interruption!
- Execute callback & remove the task
 - Execute callback **ONLY** after time out
 - Remove the task from array after executing callback

Time point: INF	Time point: 40 callback	Time point: 30 callback
-----------------	----------------------------	----------------------------

Questions?